

# Final Assessment 1

Kaviyadevi(20106064)

```
In [1]: #importing libraries
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
```

```
In [2]: #importing dataset
data1=pd.read_csv(r"C:\Users\user\Downloads\madrid_2010.csv")
data1
```

Out[2]:

	date	BEN	CO	EBE	MXV	NMHC	NO <sub>2</sub>	NO <sub>x</sub>	OXY	O <sub>3</sub>	PM <sub>10</sub>
0	2010-03-01 01:00:00	NaN	0.29	NaN	NaN	NaN	25.090000	29.219999	NaN	68.930000	NaN
1	2010-03-01 01:00:00	NaN	0.27	NaN	NaN	NaN	24.879999	30.040001	NaN	NaN	NaN
2	2010-03-01 01:00:00	NaN	0.28	NaN	NaN	NaN	17.410000	20.540001	NaN	72.120003	NaN
3	2010-03-01 01:00:00	0.38	0.24	1.74	NaN	0.05	15.610000	21.080000	NaN	72.970001	19.4100
4	2010-03-01 01:00:00	0.79	NaN	1.32	NaN	NaN	21.430000	26.070000	NaN	NaN	24.6700
...	...	...	...	...	...	...	...	...	...	...	...
209443	2010-08-01 00:00:00	NaN	0.55	NaN	NaN	NaN	125.000000	219.899994	NaN	25.379999	NaN
209444	2010-08-01 00:00:00	NaN	0.27	NaN	NaN	NaN	45.709999	47.410000	NaN	NaN	51.2599
209445	2010-08-01 00:00:00	NaN	NaN	NaN	NaN	0.24	46.560001	49.040001	NaN	46.250000	NaN
209446	2010-08-01 00:00:00	NaN	NaN	NaN	NaN	NaN	46.770000	50.119999	NaN	77.709999	NaN
209447	2010-08-01 00:00:00	0.92	0.43	0.71	NaN	0.25	76.330002	88.190002	NaN	52.259998	47.1500

209448 rows × 12 columns

In [3]: data1.info()

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 209448 entries, 0 to 209447
Data columns (total 17 columns):
#   Column      Non-Null Count  Dtype
---  -
0   date        209448 non-null  object
1   BEN         60268 non-null   float64
2   CO          94982 non-null   float64
3   EBE         60253 non-null   float64
4   MXY         6750 non-null    float64
5   NMHC        51727 non-null   float64
6   NO_2        208219 non-null   float64
7   NOx         208210 non-null   float64
8   OXY         6750 non-null    float64
9   O_3         126684 non-null   float64
10  PM10        106186 non-null   float64
11  PM25        55514 non-null   float64
12  PXY         6740 non-null    float64
13  SO_2        93184 non-null   float64
14  TCH         51730 non-null   float64
15  TOL         60171 non-null   float64
16  station     209448 non-null   int64
dtypes: float64(15), int64(1), object(1)
memory usage: 27.2+ MB
```

In [4]: data=data1.head(50000)

```
In [5]: #filling null values
df=data.fillna(0)
df
```

Out[5]:

	N	CO	EBE	MXV	NMHC	NO_2	NOx	OXY	O_3	PM10	PM25	PXY	SO_2
0	0.29	0.00	0.0	0.00	25.090000	29.219999	0.0	68.930000	0.000000	0.000000	0.0	10.1	
1	0.27	0.00	0.0	0.00	24.879999	30.040001	0.0	0.000000	0.000000	0.000000	0.0	12.2	
2	0.28	0.00	0.0	0.00	17.410000	20.540001	0.0	72.120003	0.000000	0.000000	0.0	0.0	
3	0.24	1.74	0.0	0.05	15.610000	21.080000	0.0	72.970001	19.410000	7.870000	0.0	10.0	
4	0.00	1.32	0.0	0.00	21.430000	26.070000	0.0	0.000000	24.670000	22.030001	0.0	10.6	
...	...	...	...	...	...	...	...	...	...	...	...	...	
11	0.23	1.02	0.0	0.12	32.910000	38.000000	0.0	57.400002	24.389999	13.210000	0.0	5.5	
10	0.00	0.23	0.0	0.00	21.629999	25.700001	0.0	0.000000	17.719999	10.100000	0.0	8.1	
3	0.00	0.25	0.0	0.00	17.030001	21.040001	0.0	0.000000	0.000000	0.000000	0.0	0.0	
10	0.00	0.00	0.0	0.00	28.639999	30.980000	0.0	0.000000	20.400000	0.000000	0.0	6.5	
10	0.22	0.00	0.0	0.00	55.360001	63.799999	0.0	54.169998	0.000000	0.000000	0.0	0.0	

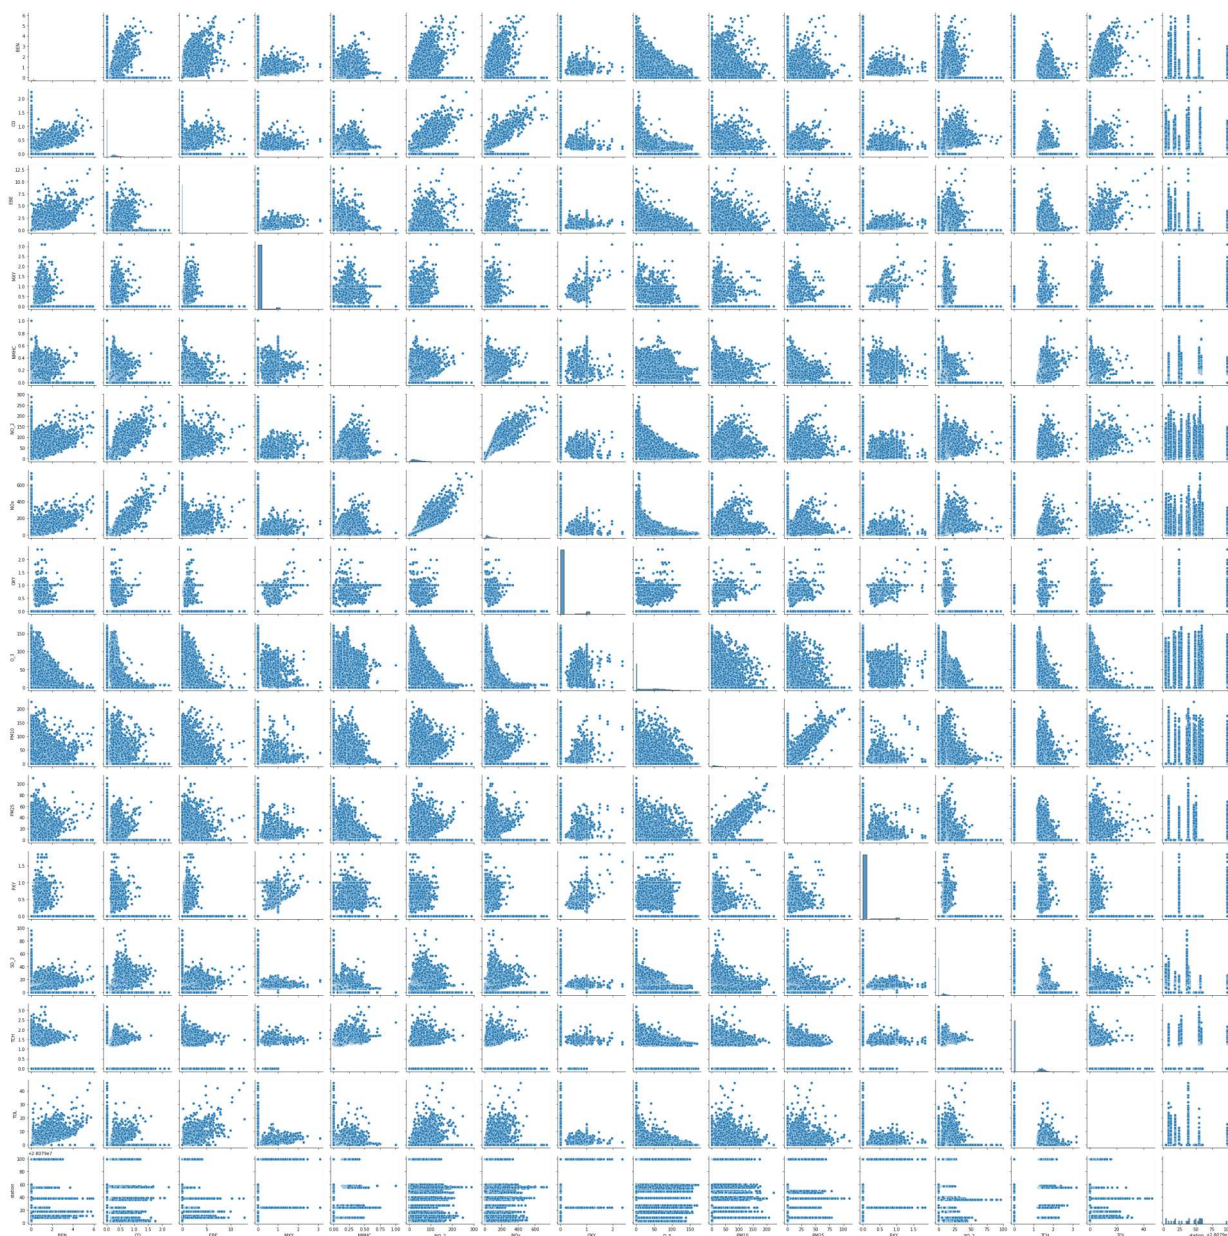
12mns

```
In [6]: df.columns
```

Out[6]: Index(['date', 'BEN', 'CO', 'EBE', 'MXV', 'NMHC', 'NO\_2', 'NOx', 'OXY', 'O\_3', 'PM10', 'PM25', 'PXY', 'SO\_2', 'TCH', 'TOL', 'station'], dtype='object')

```
In [7]: sns.pairplot(df)
```

```
Out[7]: <seaborn.axisgrid.PairGrid at 0x1b1f57f78b0>
```




```
In [9]: sns.distplot(data["NOx"])
```

C:\ProgramData\Anaconda3\lib\site-packages\seaborn\distributions.py:2557: FutureWarning: `distplot` is a deprecated function and will be removed in a future version. Please adapt your code to use either `displot` (a figure-level function with similar flexibility) or `histplot` (an axes-level function for histograms).

```
warnings.warn(msg, FutureWarning)
```

```
Out[9]: <AxesSubplot:xlabel='NOx', ylabel='Density'>
```



## MODEL BUILDING

### 1.Linear Regression

```
In [14]: df1=df[['BEN', 'CO', 'EBE', 'MX', 'NMHC', 'NO_2', 'NOx', 'OXY', 'O_3',  
               'PM10', 'PXY', 'SO_2', 'TCH', 'TOL', 'station']]
```

```
In [15]: x=df1[['BEN', 'CO', 'EBE', 'MX', 'NMHC', 'NO_2', 'OXY', 'O_3', 'PM10', 'PXY', 'SO_2', 'TCH', 'TOL', 'station']]  
y=df1[['NOx']]
```

```
In [16]: #split the dataset into training and test
from sklearn.model_selection import train_test_split
x_train,x_test,y_train,y_test=train_test_split(x,y,test_size=0.3)
```

```
In [17]: from sklearn.linear_model import LinearRegression
lr=LinearRegression()
lr.fit(x_train,y_train)
```


Out[17]: LinearRegression()

```
In [18]: print(lr.intercept_)

[-1365705.55708257]
```

```
In [19]: prediction = lr.predict(x_test)
plt.scatter(y_test,prediction)
```

Out[19]: <matplotlib.collections.PathCollection at 0x1b19d063910>



```
In [20]: print(lr.score(x_test,y_test))

0.8692089018915892
```

## 2.Ridge Regression

```
In [21]: from sklearn.linear_model import Ridge
```

```
In [22]: rr=Ridge(alpha=10)
rr.fit(x_train,y_train)
```

Out[22]: Ridge(alpha=10)

```
In [23]: rr.score(x_test,y_test)
```

Out[23]: 0.8691959334056143

## 3.Lasso Regression

```
In [24]: from sklearn.linear_model import Lasso
```

```
In [25]: la=Lasso(alpha=10)
la.fit(x_train,y_train)
```

```
Out[25]: Lasso(alpha=10)
```

```
In [26]: la.score(x_test,y_test)
```

```
Out[26]: 0.8604972643160442
```

## 4.ElasticNet Regression

```
In [27]: from sklearn.linear_model import ElasticNet
en=ElasticNet()
en.fit(x_train,y_train)
```

```
Out[27]: ElasticNet()
```

```
In [28]: print(en.coef_)
```

```
[ 0.00000000e+00  4.78422177e-02  2.22102267e-01 -0.00000000e+00
 0.00000000e+00  1.70494634e+00 -0.00000000e+00  8.61202032e-03
-4.04713928e-02 -0.00000000e+00  5.35160842e-01  0.00000000e+00
 1.58818980e-01  1.57322777e-03]
```

```
In [29]: print(en.predict(x_test))
```

```
[ 42.53741104 124.67896895 -13.2848226 ... 37.29012965 20.50131458
 99.52221428]
```

```
In [30]: print(en.score(x_test,y_test))
```

```
0.8619496362026571
```

## 5.Logistic Regression

```
In [31]: from sklearn.linear_model import LogisticRegression
```

```
In [32]: feature_matrix = df1.iloc[:,0:16]
target_vector = df1.iloc[:,-1]
```

```
In [33]: feature_matrix.shape
```

```
Out[33]: (50000, 15)
```

```
In [34]: target_vector.shape
```

```
Out[34]: (50000,)
```

```
In [35]: from sklearn.preprocessing import StandardScaler
```

```
In [36]: fs=StandardScaler().fit_transform(feature_matrix)
```

```
In [37]: logr = LogisticRegression()  
logr.fit(fs,target_vector)
```

C:\ProgramData\Anaconda3\lib\site-packages\sklearn\linear\_model\\_logistic.py:763: ConvergenceWarning: lbfgs failed to converge (status=1):  
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.

Increase the number of iterations (max\_iter) or scale the data as shown in:  
<https://scikit-learn.org/stable/modules/preprocessing.html> (<https://scikit-learn.org/stable/modules/preprocessing.html>)  
Please also refer to the documentation for alternative solver options:  
[https://scikit-learn.org/stable/modules/linear\\_model.html#logistic-regression](https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression) ([https://scikit-learn.org/stable/modules/linear\\_model.html#logistic-regression](https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression))  
n\_iter\_i = \_check\_optimize\_result(

```
Out[37]: LogisticRegression()
```

```
In [38]: observation=[[1,2,3,4,5,6,7,8,9,10,11,12,13,14,15]]
```

```
In [39]: prediction=logr.predict(observation)  
print(prediction)
```

```
[28079099]
```

```
In [40]: logr.classes_
```

```
Out[40]: array([28079003, 28079004, 28079008, 28079011, 28079016, 28079017,  
                28079018, 28079024, 28079027, 28079036, 28079038, 28079039,  
                28079040, 28079047, 28079048, 28079049, 28079050, 28079054,  
                28079055, 28079056, 28079057, 28079058, 28079059, 28079060,  
                28079099], dtype=int64)
```

```
In [41]: logr.score(fs,target_vector)
```

```
Out[41]: 0.98068
```

## 6.Random Forest



```
In [56]: df1=df[['BEN', 'CO', 'EBE', 'MXY', 'NMHC', 'NO_2', 'NOx', 'OXY', 'O_3', 'PM10', 'P  
x=df1[['BEN', 'CO', 'EBE', 'NMHC', 'NO_2', 'OXY', 'PM10', 'PXY', 'SO_2', 'TCH', 'T  
y=df['station']
```

```
In [57]: from sklearn.model_selection import train_test_split  
x_train,x_test,y_train,y_test=train_test_split(x,y,test_size=45)
```

```
In [58]: from sklearn.ensemble import RandomForestClassifier  
rfc = RandomForestClassifier()  
rfc.fit(x_train,y_train)
```

Out[58]: RandomForestClassifier()

```
In [59]: parameters = {'max_depth':[1,2,3,4,5],  
                        'min_samples_leaf':[5,10,15,20,25],  
                        'n_estimators':[10,20,30,40,50]}
```

```
In [60]: from sklearn.model_selection import GridSearchCV  
  
grid_search = GridSearchCV(estimator=rfc,param_grid=parameters,cv=2,scoring='acc  
grid_search.fit(x_train,y_train)
```

Out[60]: GridSearchCV(cv=2, estimator=RandomForestClassifier(),  
param\_grid={'max\_depth': [1, 2, 3, 4, 5],  
'min\_samples\_leaf': [5, 10, 15, 20, 25],  
'n\_estimators': [10, 20, 30, 40, 50]},  
scoring='accuracy')

```
In [61]: grid_search.best_score_
```

Out[61]: 0.9550593720536289

```
In [62]: rfc_best = grid_search.best_estimator_
```

