

Final Assessment 1

Kaviyadevi(20106064)

```
In [1]: #importing Libraries  
import numpy as np  
import pandas as pd  
import matplotlib.pyplot as plt  
import seaborn as sns
```

```
In [2]: #importing dataset
data1=pd.read_csv(r"C:\Users\user\Downloads\madrid_2012.csv")
data1
```

Out[2]:

	date	BEN	CO	EBE	NMHC	NO	NO_2	O_3	PM10	PM25	SO_2	TCH	TOL	s
0	2012-09-01 01:00:00	NaN	0.2	NaN	NaN	7.0	18.0	NaN	NaN	NaN	2.0	NaN	NaN	280
1	2012-09-01 01:00:00	0.3	0.3	0.7	NaN	3.0	18.0	55.0	10.0	9.0	1.0	NaN	2.4	280
2	2012-09-01 01:00:00	0.4	NaN	0.7	NaN	2.0	10.0	NaN	NaN	NaN	NaN	NaN	1.5	280
3	2012-09-01 01:00:00	NaN	0.2	NaN	NaN	1.0	6.0	50.0	NaN	NaN	NaN	NaN	NaN	280
4	2012-09-01 01:00:00	NaN	NaN	NaN	NaN	1.0	13.0	54.0	NaN	NaN	3.0	NaN	NaN	280
...
210715	2012-03-01 00:00:00	NaN	0.6	NaN	NaN	37.0	84.0	14.0	NaN	NaN	NaN	NaN	NaN	280
210716	2012-03-01 00:00:00	NaN	0.4	NaN	NaN	5.0	76.0	NaN	17.0	NaN	7.0	NaN	NaN	280
210717	2012-03-01 00:00:00	NaN	NaN	NaN	0.34	3.0	41.0	24.0	NaN	NaN	NaN	1.34	NaN	280
210718	2012-03-01 00:00:00	NaN	NaN	NaN	NaN	2.0	44.0	36.0	NaN	NaN	NaN	NaN	NaN	280
210719	2012-03-01 00:00:00	NaN	NaN	NaN	NaN	2.0	56.0	40.0	18.0	NaN	NaN	NaN	NaN	280

210720 rows × 14 columns



```
In [3]: data1.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 210720 entries, 0 to 210719
Data columns (total 14 columns):
#   Column      Non-Null Count  Dtype
---  -
0   date        210720 non-null  object
1   BEN         51511 non-null   float64
2   CO          87097 non-null   float64
3   EBE         51482 non-null   float64
4   NMHC        30736 non-null   float64
5   NO          209871 non-null  float64
6   NO_2        209872 non-null  float64
7   O_3         122339 non-null  float64
8   PM10        104838 non-null  float64
9   PM25        52164 non-null   float64
10  SO_2        87333 non-null   float64
11  TCH         30736 non-null   float64
12  TOL         51373 non-null   float64
13  station     210720 non-null  int64
dtypes: float64(12), int64(1), object(1)
memory usage: 22.5+ MB
```

```
In [4]: data=data1.head(50000)
```

```
In [5]: #filling null values
df=data.fillna(0)
df
```

Out[5]:

	date	BEN	CO	EBE	NMHC	NO	NO_2	O_3	PM10	PM25	SO_2	TCH	TOL	station
0	2012-09-01 01:00:00	0.0	0.2	0.0	0.00	7.0	18.0	0.0	0.0	0.0	2.0	0.00	0.0	280790
1	2012-09-01 01:00:00	0.3	0.3	0.7	0.00	3.0	18.0	55.0	10.0	9.0	1.0	0.00	2.4	280790
2	2012-09-01 01:00:00	0.4	0.0	0.7	0.00	2.0	10.0	0.0	0.0	0.0	0.0	0.00	1.5	280790
3	2012-09-01 01:00:00	0.0	0.2	0.0	0.00	1.0	6.0	50.0	0.0	0.0	0.0	0.00	0.0	280790
4	2012-09-01 01:00:00	0.0	0.0	0.0	0.00	1.0	13.0	54.0	0.0	0.0	3.0	0.00	0.0	280790
...
49995	2012-03-26 20:00:00	0.0	0.3	0.0	0.00	7.0	54.0	33.0	0.0	0.0	0.0	0.00	0.0	280790
49996	2012-03-26 20:00:00	0.0	0.0	0.0	0.00	3.0	29.0	62.0	0.0	0.0	3.0	0.00	0.0	280790
49997	2012-03-26 20:00:00	0.2	0.4	1.0	0.00	4.0	36.0	62.0	29.0	0.0	2.0	0.00	0.4	280790
49998	2012-03-26 20:00:00	0.3	0.2	0.2	0.17	1.0	26.0	72.0	25.0	14.0	2.0	1.23	0.8	280790
49999	2012-03-26 20:00:00	0.0	0.0	0.0	0.17	1.0	31.0	67.0	0.0	0.0	0.0	1.22	0.0	280790

50000 rows × 14 columns

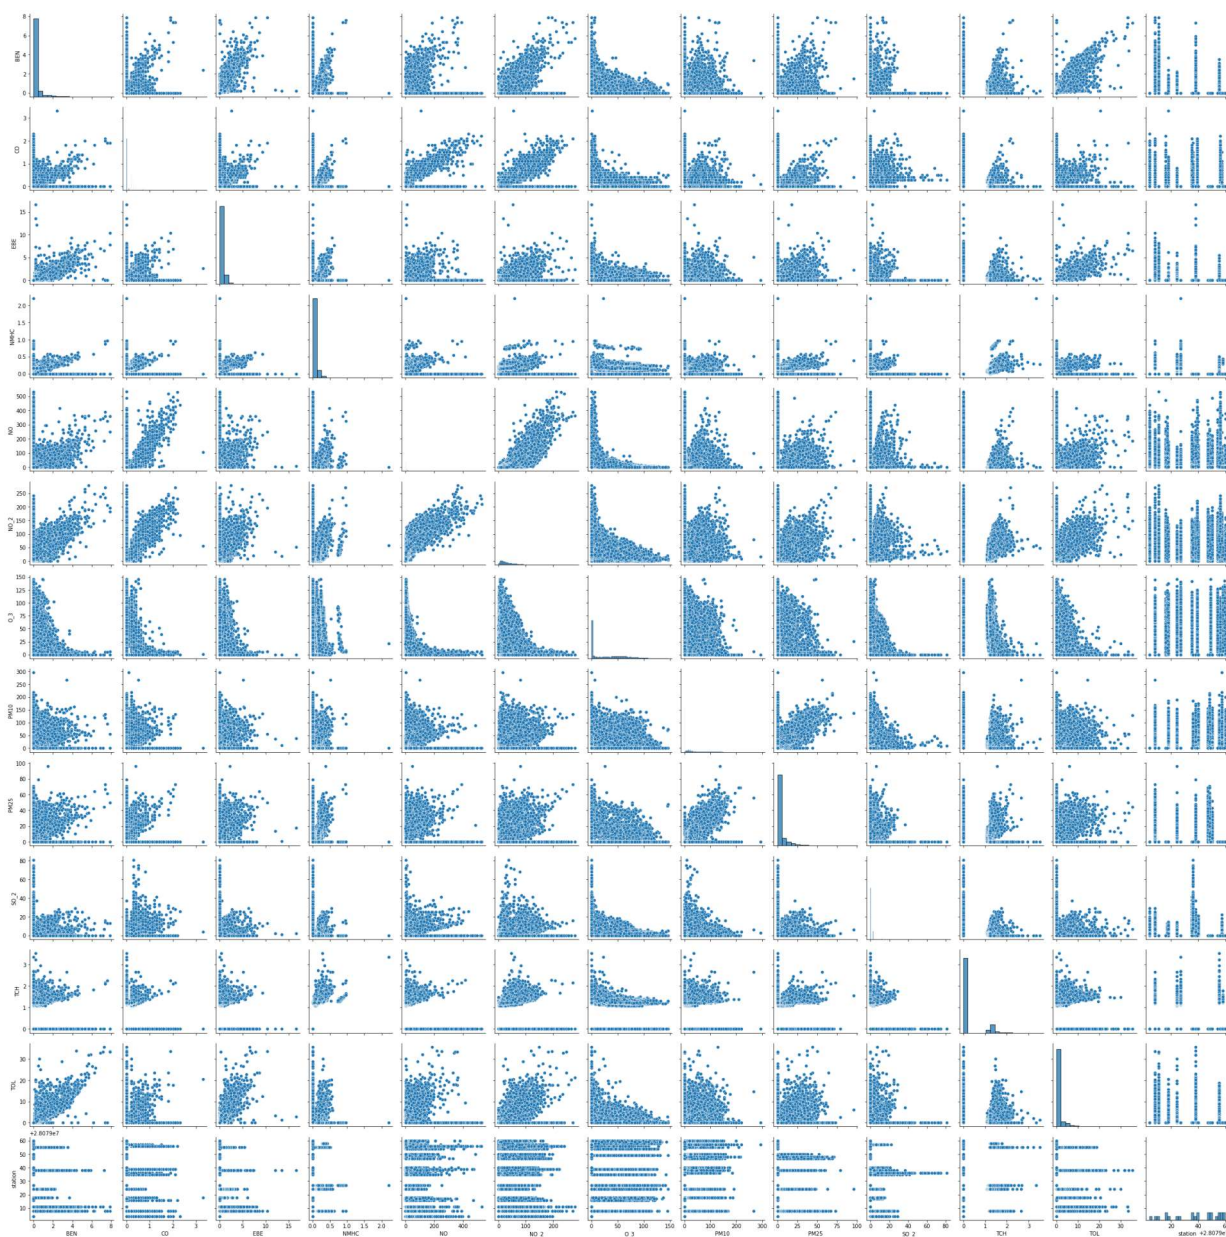


```
In [6]: df.columns
```

Out[6]: Index(['date', 'BEN', 'CO', 'EBE', 'NMHC', 'NO', 'NO_2', 'O_3', 'PM10', 'PM25', 'SO_2', 'TCH', 'TOL', 'station'], dtype='object')

```
In [7]: sns.pairplot(df)
```

```
Out[7]: <seaborn.axisgrid.PairGrid at 0x1d115d93220>
```




```
In [8]: sns.distplot(data["BEN"])
```

C:\ProgramData\Anaconda3\lib\site-packages\seaborn\distributions.py:2557: FutureWarning: `distplot` is a deprecated function and will be removed in a future version. Please adapt your code to use either `displot` (a figure-level function with similar flexibility) or `histplot` (an axes-level function for histograms).

```
warnings.warn(msg, FutureWarning)
```

```
Out[8]: <AxesSubplot:xlabel='BEN', ylabel='Density'>
```



MODEL BUILDING

1.Linear Regression

```
In [9]: df1=df[['BEN', 'CO', 'EBE', 'NMHC', 'NO', 'NO_2', 'O_3', 'PM10', 'PM25',  
              'SO_2', 'TCH', 'TOL', 'station']]
```

```
In [10]: x=df1[['CO', 'EBE', 'NMHC', 'NO', 'NO_2', 'O_3', 'PM10', 'PM25',  
              'SO_2', 'TCH', 'TOL', 'station']]  
y=df1[['BEN']]
```

```
In [11]: #split the dataset into training and test  
from sklearn.model_selection import train_test_split  
x_train,x_test,y_train,y_test=train_test_split(x,y,test_size=0.3)
```


```
In [12]: from sklearn.linear_model import LinearRegression  
lr=LinearRegression()  
lr.fit(x_train,y_train)
```

Out[12]: LinearRegression()

```
In [13]: print(lr.intercept_)  
  
[57987.49167714]
```

```
In [14]: prediction = lr.predict(x_test)  
plt.scatter(y_test,prediction)
```

Out[14]: <matplotlib.collections.PathCollection at 0x1d12cc75eb0>



```
In [15]: print(lr.score(x_test,y_test))  
  
0.7537693237620529
```

2.Ridge Regression

```
In [16]: from sklearn.linear_model import Ridge
```

```
In [17]: rr=Ridge(alpha=10)  
rr.fit(x_train,y_train)
```

Out[17]: Ridge(alpha=10)

```
In [18]: rr.score(x_test,y_test)
```

```
Out[18]: 0.7534104823730857
```

3.Lasso Regression

```
In [19]: from sklearn.linear_model import Lasso
```

```
In [20]: la=Lasso(alpha=10)
la.fit(x_train,y_train)
```

```
Out[20]: Lasso(alpha=10)
```

```
In [21]: la.score(x_test,y_test)
```

```
Out[21]: -2.64929742690434e-05
```

4.ElasticNet Regression

```
In [22]: from sklearn.linear_model import ElasticNet
en=ElasticNet()
en.fit(x_train,y_train)
```

```
Out[22]: ElasticNet()
```

```
In [23]: print(en.coef_)
```

```
[-0.          0.          0.          0.00053264  0.00209504 -0.
 0.00123358  0.          -0.          0.          0.04399504 -0.00614486]
```

```
In [24]: print(en.predict(x_test))
```

```
[0.05702859 0.03235718 0.10240988 ... 0.08483021 0.00349833 0.07321679]
```

```
In [25]: print(en.score(x_test,y_test))
```

```
0.40344157128662583
```

5.Logistic Regression

```
In [26]: from sklearn.linear_model import LogisticRegression
```

```
In [27]: feature_matrix = df1.iloc[:,0:14]
target_vector = df1.iloc[:,-1]
```



```
In [28]: feature_matrix.shape
```

```
Out[28]: (50000, 11)
```

```
In [29]: target_vector.shape
```

```
Out[29]: (50000,)
```

```
In [30]: from sklearn.preprocessing import StandardScaler
```

```
In [31]: fs=StandardScaler().fit_transform(feature_matrix)
```

```
In [ ]: logr = LogisticRegression()  
logr.fit(fs,target_vector)
```

```
In [ ]: observation=[[1,2,3,4,5,6,7,8,9,10,11,12,13,14]]
```

```
In [ ]: prediction=logr.predict(observation)  
print(prediction)
```

```
In [ ]: logr.classes_
```

```
In [ ]: logr.score(fs,target_vector)
```

6.Random Forest

```
In [ ]: df1=df[['BEN', 'CO', 'EBE', 'NMHC', 'NO', 'NO_2', 'O_3', 'PM10', 'PM25',  
              'SO_2', 'TCH', 'TOL', 'station']]  
x=df1[['CO', 'EBE', 'NMHC', 'NO', 'NO_2', 'O_3', 'PM10', 'PM25',  
       'SO_2', 'TCH', 'TOL']]  
y=df1['station']
```

```
In [ ]: from sklearn.model_selection import train_test_split  
x_train,x_test,y_train,y_test=train_test_split(x,y,test_size=0.3)
```

```
In [ ]: from sklearn.ensemble import RandomForestClassifier  
rfc = RandomForestClassifier()  
rfc.fit(x_train,y_train)
```

```
In [ ]: parameters = {'max_depth':[1,2,3,4,5],  
                      'min_samples_leaf':[5,10,15,20,25],  
                      'n_estimators':[10,20,30,40,50]}
```

```
In [ ]: from sklearn.model_selection import GridSearchCV

grid_search = GridSearchCV(estimator=rfc,param_grid=parameters,cv=2,scoring='acc
grid_search.fit(x_train,y_train)
```

```
In [ ]: grid_search.best_score_
```

```
In [ ]: rfc_best = grid_search.best_estimator_
```

```
In [ ]: from sklearn.tree import plot_tree

plt.figure(figsize=(80,40))
plot_tree(rfc_best.estimators_[5],feature_names=x.columns,filled=True)
```

Results

```
1.Linear regression : 0.8455591055403516
2.Ridge regression : 0.845561600678548
3.Lasso regression : -0.00022865063965538113
4.Elasticnet regression : 0.6309166326143854
5.Logistic regresssion : 0.75178
6.Random forest regression : 0.7791428571428571

Hence Linear regression gives high accuracy for the madrid_2012 model.
```