

# Final Assessment 1

Kaviyadevi(20106064)

```
In [1]: #importing Libraries  
import numpy as np  
import pandas as pd  
import matplotlib.pyplot as plt  
import seaborn as sns
```

```
In [2]: #importing dataset
data1=pd.read_csv(r"C:\Users\user\Downloads\madrid_2018.csv")
data1
```

Out[2]:

	date	BEN	CH4	CO	EBE	NMHC	NO	NO_2	NOx	O_3	PM10	PM25	SO_2	TC
0	2018-03-01 01:00:00	NaN	NaN	0.3	NaN	NaN	1.0	29.0	31.0	NaN	NaN	NaN	2.0	NaN
1	2018-03-01 01:00:00	0.5	1.39	0.3	0.2	0.02	6.0	40.0	49.0	52.0	5.0	4.0	3.0	1.4
2	2018-03-01 01:00:00	0.4	NaN	NaN	0.2	NaN	4.0	41.0	47.0	NaN	NaN	NaN	NaN	NaN
3	2018-03-01 01:00:00	NaN	NaN	0.3	NaN	NaN	1.0	35.0	37.0	54.0	NaN	NaN	NaN	NaN
4	2018-03-01 01:00:00	NaN	NaN	NaN	NaN	NaN	1.0	27.0	29.0	49.0	NaN	NaN	3.0	NaN
...	...	...	...	...	...	...	...	...	...	...	...	...	...	...
69091	2018-02-01 00:00:00	NaN	NaN	0.5	NaN	NaN	66.0	91.0	192.0	1.0	35.0	22.0	NaN	NaN
69092	2018-02-01 00:00:00	NaN	NaN	0.7	NaN	NaN	87.0	107.0	241.0	NaN	29.0	NaN	15.0	NaN
69093	2018-02-01 00:00:00	NaN	NaN	NaN	NaN	NaN	28.0	48.0	91.0	2.0	NaN	NaN	NaN	NaN
69094	2018-02-01 00:00:00	NaN	NaN	NaN	NaN	NaN	141.0	103.0	320.0	2.0	NaN	NaN	NaN	NaN
69095	2018-02-01 00:00:00	NaN	NaN	NaN	NaN	NaN	69.0	96.0	202.0	3.0	26.0	NaN	NaN	NaN

69096 rows × 16 columns

In [3]: data1.info()

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 69096 entries, 0 to 69095
Data columns (total 16 columns):
#   Column      Non-Null Count  Dtype
---  -
0   date        69096 non-null  object
1   BEN         16950 non-null  float64
2   CH4         8440 non-null   float64
3   CO          28598 non-null  float64
4   EBE         16949 non-null  float64
5   NMHC        8440 non-null   float64
6   NO          68826 non-null  float64
7   NO_2        68826 non-null  float64
8   NOx         68826 non-null  float64
9   O_3         40049 non-null  float64
10  PM10        36911 non-null  float64
11  PM25        18912 non-null  float64
12  SO_2        28586 non-null  float64
13  TCH         8440 non-null   float64
14  TOL         16950 non-null  float64
15  station     69096 non-null  int64
dtypes: float64(14), int64(1), object(1)
memory usage: 8.4+ MB
```


In [4]: data=data1.head(50000)

```
In [5]: #filling null values
df=data.fillna(0)
df
```

Out[5]:

	date	BEN	CH4	CO	EBE	NMHC	NO	NO_2	NOx	O_3	PM10	PM25	SO_2	TCH
0	2018-03-01 01:00:00	0.0	0.00	0.3	0.0	0.00	1.0	29.0	31.0	0.0	0.0	0.0	2.0	0.00
1	2018-03-01 01:00:00	0.5	1.39	0.3	0.2	0.02	6.0	40.0	49.0	52.0	5.0	4.0	3.0	1.41
2	2018-03-01 01:00:00	0.4	0.00	0.0	0.2	0.00	4.0	41.0	47.0	0.0	0.0	0.0	0.0	0.00
3	2018-03-01 01:00:00	0.0	0.00	0.3	0.0	0.00	1.0	35.0	37.0	54.0	0.0	0.0	0.0	0.00
4	2018-03-01 01:00:00	0.0	0.00	0.0	0.0	0.00	1.0	27.0	29.0	49.0	0.0	0.0	3.0	0.00
...	...	...	...	...	...	...	...	...	...	...	...	...	...	...
49995	2018-02-26 21:00:00	0.0	0.00	0.6	0.0	0.00	18.0	109.0	137.0	11.0	0.0	0.0	0.0	0.00
49996	2018-02-26 21:00:00	0.0	0.00	0.0	0.0	0.00	171.0	160.0	422.0	1.0	0.0	0.0	12.0	0.00
49997	2018-02-26 21:00:00	0.8	0.00	0.8	0.7	0.00	19.0	109.0	138.0	10.0	25.0	0.0	7.0	0.00
49998	2018-02-26 21:00:00	0.7	1.10	0.3	0.4	0.09	1.0	81.0	82.0	20.0	20.0	11.0	4.0	1.19
49999	2018-02-26 21:00:00	0.0	0.00	0.0	0.0	0.00	30.0	113.0	160.0	8.0	0.0	0.0	0.0	0.00

50000 rows × 16 columns

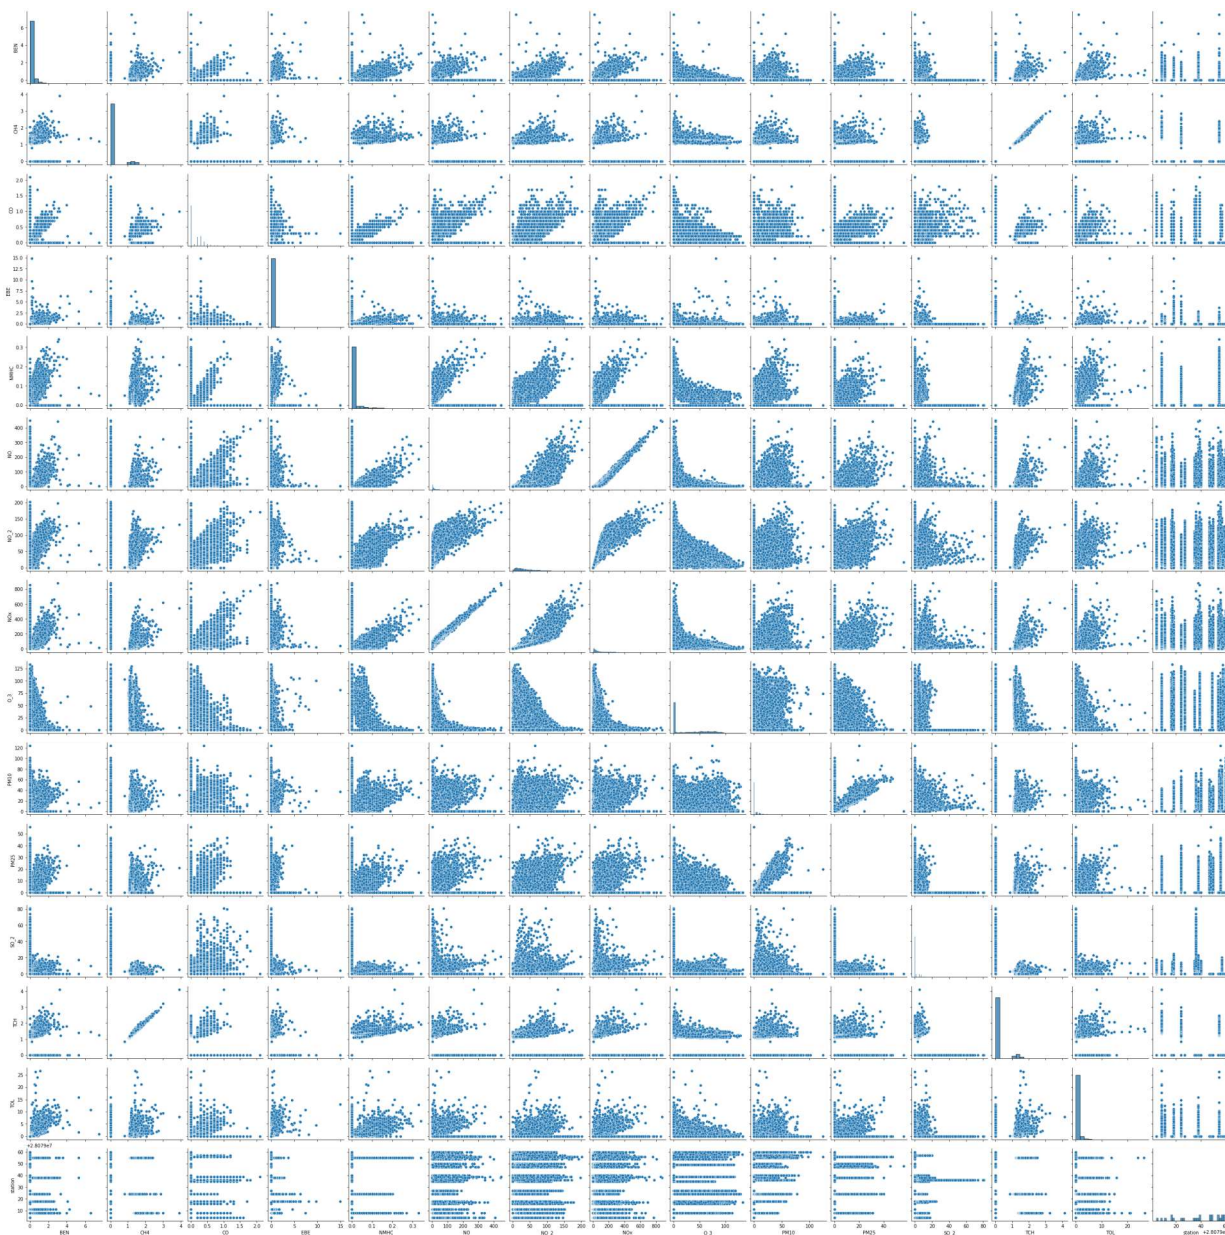


```
In [6]: df.columns
```

Out[6]: Index(['date', 'BEN', 'CH4', 'CO', 'EBE', 'NMHC', 'NO', 'NO\_2', 'NOx', 'O\_3', 'PM10', 'PM25', 'SO\_2', 'TCH', 'TOL', 'station'], dtype='object')

```
In [7]: sns.pairplot(df)
```

```
Out[7]: <seaborn.axisgrid.PairGrid at 0x21c0bdfd310>
```




```
In [8]: sns.distplot(data['NMHC'])
```

C:\ProgramData\Anaconda3\lib\site-packages\seaborn\distributions.py:2557: FutureWarning: `distplot` is a deprecated function and will be removed in a future version. Please adapt your code to use either `displot` (a figure-level function with similar flexibility) or `histplot` (an axes-level function for histograms).

```
warnings.warn(msg, FutureWarning)
```

```
Out[8]: <AxesSubplot:xlabel='NMHC', ylabel='Density'>
```



## MODEL BUILDING

### 1.Linear Regression

```
In [9]: df1=df[['BEN', 'CO', 'EBE', 'NMHC', 'NO', 'NO_2', 'O_3', 'PM10', 'PM25', 'SO_2',
```

```
In [10]: x=df1[['BEN', 'CO', 'EBE', 'NO', 'NO_2', 'O_3', 'PM10', 'PM25', 'SO_2', 'TCH',  
y=df1[['NMHC']]
```

```
In [11]: #split the dataset into training and test  
from sklearn.model_selection import train_test_split  
x_train,x_test,y_train,y_test=train_test_split(x,y,test_size=0.3)
```


```
In [12]: from sklearn.linear_model import LinearRegression  
lr=LinearRegression()  
lr.fit(x_train,y_train)
```

Out[12]: LinearRegression()

```
In [13]: print(lr.intercept_)  
  
[-5546.12509531]
```

```
In [14]: prediction = lr.predict(x_test)  
plt.scatter(y_test,prediction)
```

Out[14]: <matplotlib.collections.PathCollection at 0x21c28772c40>



```
In [15]: print(lr.score(x_test,y_test))  
  
0.7850736800457832
```

## 2.Ridge Regression

```
In [16]: from sklearn.linear_model import Ridge
```

```
In [17]: rr=Ridge(alpha=10)  
rr.fit(x_train,y_train)
```

Out[17]: Ridge(alpha=10)

```
In [18]: rr.score(x_test,y_test)
```

```
Out[18]: 0.7851046395472827
```

## 3.Lasso Regression

```
In [19]: from sklearn.linear_model import Lasso
```

```
In [20]: la=Lasso(alpha=10)
la.fit(x_train,y_train)
```

```
Out[20]: Lasso(alpha=10)
```

```
In [21]: la.score(x_test,y_test)
```

```
Out[21]: -2.5541386561300783e-08
```

## 4.ElasticNet Regression

```
In [22]: from sklearn.linear_model import ElasticNet
en=ElasticNet()
en.fit(x_train,y_train)
```

```
Out[22]: ElasticNet()
```

```
In [23]: print(en.coef_)
```

```
[ 0.  0.  0.  0.  0. -0.  0.  0.  0.  0.  0. -0.]
```

```
In [24]: print(en.predict(x_test))
```

```
[0.00810543 0.00810543 0.00810543 ... 0.00810543 0.00810543 0.00810543]
```

```
In [25]: print(en.score(x_test,y_test))
```

```
-2.5541386561300783e-08
```

## 5.Logistic Regression

```
In [26]: from sklearn.linear_model import LogisticRegression
```

```
In [27]: feature_matrix = df1.iloc[:,0:11]
target_vector = df1.iloc[:,-1]
```



```
In [28]: feature_matrix.shape
```

```
Out[28]: (50000, 11)
```

```
In [29]: target_vector.shape
```

```
Out[29]: (50000,)
```

```
In [30]: from sklearn.preprocessing import StandardScaler
```

```
In [31]: fs=StandardScaler().fit_transform(feature_matrix)
```

```
In [32]: logr = LogisticRegression()  
logr.fit(fs,target_vector)
```

C:\ProgramData\Anaconda3\lib\site-packages\sklearn\linear\_model\\_logistic.py:763: ConvergenceWarning: lbfgs failed to converge (status=1):  
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.

Increase the number of iterations (max\_iter) or scale the data as shown in:  
<https://scikit-learn.org/stable/modules/preprocessing.html> (<https://scikit-learn.org/stable/modules/preprocessing.html>)  
Please also refer to the documentation for alternative solver options:  
[https://scikit-learn.org/stable/modules/linear\\_model.html#logistic-regression](https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression) ([https://scikit-learn.org/stable/modules/linear\\_model.html#logistic-regression](https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression))  
n\_iter\_i = \_check\_optimize\_result(

```
Out[32]: LogisticRegression()
```

```
In [33]: observation=[[1,2,3,4,5,6,7,8,9,10,11]]
```

```
In [34]: prediction=logr.predict(observation)  
print(prediction)
```

```
[28079008]
```

```
In [35]: logr.classes_
```

```
Out[35]: array([28079004, 28079008, 28079011, 28079016, 28079017, 28079018,  
                28079024, 28079027, 28079035, 28079036, 28079038, 28079039,  
                28079040, 28079047, 28079048, 28079049, 28079050, 28079054,  
                28079055, 28079056, 28079057, 28079058, 28079059, 28079060],  
              dtype=int64)
```

```
In [36]: logr.score(fs,target_vector)
```

```
Out[36]: 0.73
```

## 6.Random Forest

```
In [37]: df1=df[['date', 'BEN', 'CO', 'EBE', 'NMHC', 'NO', 'NO_2', 'O_3', 'PM10', 'PM25',  
x=df1[['CO', 'NMHC', 'NO_2', 'O_3', 'PM10', 'SO_2', 'TCH', 'TOL']]  
y=df1['station']
```

```
In [38]: from sklearn.model_selection import train_test_split  
x_train,x_test,y_train,y_test=train_test_split(x,y,test_size=0.3)
```

```
In [39]: from sklearn.ensemble import RandomForestClassifier  
rfc = RandomForestClassifier()  
rfc.fit(x_train,y_train)
```

Out[39]: RandomForestClassifier()

```
In [40]: parameters = {'max_depth':[1,2,3,4,5],  
                        'min_samples_leaf':[5,10,15,20,25],  
                        'n_estimators':[10,20,30,40,50]}
```

```
In [41]: from sklearn.model_selection import GridSearchCV  
  
grid_search = GridSearchCV(estimator=rfc,param_grid=parameters,cv=2,scoring='acc  
grid_search.fit(x_train,y_train)
```

Out[41]: GridSearchCV(cv=2, estimator=RandomForestClassifier(),  
param\_grid={'max\_depth': [1, 2, 3, 4, 5],  
'min\_samples\_leaf': [5, 10, 15, 20, 25],  
'n\_estimators': [10, 20, 30, 40, 50]},  
scoring='accuracy')

```
In [42]: grid_search.best_score_
```

Out[42]: 0.7129714285714286

```
In [43]: rfc_best = grid_search.best_estimator_
```

In [44]: `from sklearn.tree import plot_tree`

```
plt.figure(figsize=(80,40))
plot_tree(rfc_best.estimators_[5],feature_names=x.columns,filled=True)
0, 0, 0, 0, 0, 0, 0, 0, 1/2, 0, 0, 0, 0\n0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0] ),
  Text(3766.5, 1268.4, 'SO_2 <= 0.5\ngini = 0.857\nsamples = 6334\nvalue = [0,
1439, 0, 1461, 0, 1366, 1463, 0, 1411, 0, 0\n1501, 0, 0, 0, 0, 0, 0, 1391,
0, 0, 0, 0]'),
  Text(3627.0, 906.0, 'gini = 0.668\nsamples = 2758\nvalue = [0, 9, 0, 1461,
0, 1, 0, 0, 4, 0, 0, 1501, 0\n0, 0, 0, 0, 0, 0, 1391, 0, 0, 0, 0]'),
  Text(3906.0, 906.0, 'TOL <= 0.05\ngini = 0.75\nsamples = 3576\nvalue = [0, 1
430, 0, 0, 0, 1365, 1463, 0, 1407, 0, 0, 0\n0, 0, 0, 0, 0, 0, 0, 0, 0, 0,
0]'),
  Text(3627.0, 543.5999999999999, 'O_3 <= 69.5\ngini = 0.06\nsamples = 930\nva
lue = [0, 7, 0, 0, 0, 33, 5, 0, 1407, 0, 0, 0, 0, 0\n0, 0, 0, 0, 0, 0, 0, 0,
0, 0]'),
  Text(3487.5, 181.19999999999982, 'gini = 0.044\nsamples = 910\nvalue = [0,
7, 0, 0, 0, 24, 1, 0, 1390, 0, 0, 0, 0, 0\n0, 0, 0, 0, 0, 0, 0, 0, 0, 0]'),
  Text(3766.5, 181.19999999999982, 'gini = 0.571\nsamples = 20\nvalue = [0, 0,
0, 0, 0, 9, 4, 0, 17, 0, 0, 0, 0, 0\n0, 0, 0, 0, 0, 0, 0, 0, 0, 0]'),
  Text(4185.0, 543.5999999999999, 'TOL <= 0.25\ngini = 0.666\nsamples = 2646\n
value = [0, 1423, 0, 0, 0, 1332, 1458, 0, 0, 0, 0, 0, 0\n0, 0, 0, 0, 0, 0, 0,
0, 0, 0, 0]'),
  Text(4045.5, 181.19999999999982, 'gini = 0.311\nsamples = 513\nvalue = [0,
```

## Results

1.Linear regression : 0.7850736800457832

2.Ridge regression :0.7851046395472827

3.Lasso regression : -2.5541386561300783e-08

4.Elasticnet regression : -2.5541386561300783e-08

5.Logistic regresssion : 0.73

6.Random forest regression : 0.7129714285714286

Hence Ridge regression gives high accuracy for the madrid\_2013 model.