

Final Assessment 1

Kaviyadevi(20106064)

```
In [1]: #importing libraries  
import numpy as np  
import pandas as pd  
import matplotlib.pyplot as plt  
import seaborn as sns
```

```
In [4]: #importing dataset
data=pd.read_csv(r"C:\Users\user\Downloads\stations.csv")
data
```

Out[4]:

	id	name	address	lon	lat	elevation
0	28079004	Pza. de España	Plaza de España	-3.712247	40.423853	635
1	28079008	Escuelas Aguirre	Entre C/ Alcalá y C/ O' Donell	-3.682319	40.421564	670
2	28079011	Avda. Ramón y Cajal	Avda. Ramón y Cajal esq. C/ Príncipe de Vergara	-3.677356	40.451475	708
3	28079016	Arturo Soria	C/ Arturo Soria esq. C/ Vizconde de los Asilos	-3.639233	40.440047	693
4	28079017	Villaverde	C/. Juan Peñalver	-3.713322	40.347139	604
5	28079018	Farolillo	Calle Farolillo - C/Ervigio	-3.731853	40.394781	630
6	28079024	Casa de Campo	Casa de Campo (Terminal del Teleférico)	-3.747347	40.419356	642
7	28079027	Barajas Pueblo	C/. Júpiter, 21 (Barajas)	-3.580031	40.476928	621
8	28079035	Pza. del Carmen	Plaza del Carmen esq. Tres Cruces.	-3.703172	40.419208	659
9	28079036	Moratalaz	Avd. Moratalaz esq. Camino de los Vinateros	-3.645306	40.407947	685
10	28079038	Cuatro Caminos	Avda. Pablo Iglesias esq. C/ Marqués de Lema	-3.707128	40.445544	698
11	28079039	Barrio del Pilar	Avd. Betanzos esq. C/ Monforte de Lemos	-3.711542	40.478228	674
12	28079040	Vallecas	C/ Arroyo del Olivar esq. C/ Río Grande.	-3.651522	40.388153	677
13	28079047	Mendez Alvaro	C/ Juan de Mariana / Pza. Amanecer Mendez Alvaro	-3.686825	40.398114	599
14	28079048	Castellana	C/ Jose Gutierrez Abascal	-3.690367	40.439897	676
15	28079049	Parque del Retiro	Paseo Venezuela- Casa de Vacas	-3.682583	40.414444	662
16	28079050	Plaza Castilla	Plaza Castilla (Canal)	-3.688769	40.465572	728
17	28079054	Ensanche de Vallecas	Avda La Gavia / Avda. Las Suertes	-3.612117	40.372933	627
18	28079055	Urb. Embajada	C/ Riaño (Barajas)	-3.580747	40.462531	618
19	28079056	Pza. Fernández Ladreda	Pza. Fernández Ladreda - Avda. Oporto	-3.718728	40.384964	604
20	28079057	Sanchinarro	C/ Princesa de Eboli esq C/ Maria Tudor	-3.660503	40.494208	700
21	28079058	El Pardo	Avda. La Guardia	-3.774611	40.518058	615
22	28079059	Juan Carlos I	Parque Juan Carlos I (frente oficinas mantenim...	-3.609072	40.465250	660
23	28079060	Tres Olivos	Plaza Tres Olivos	-3.689761	40.500589	715

In [6]: data.info()

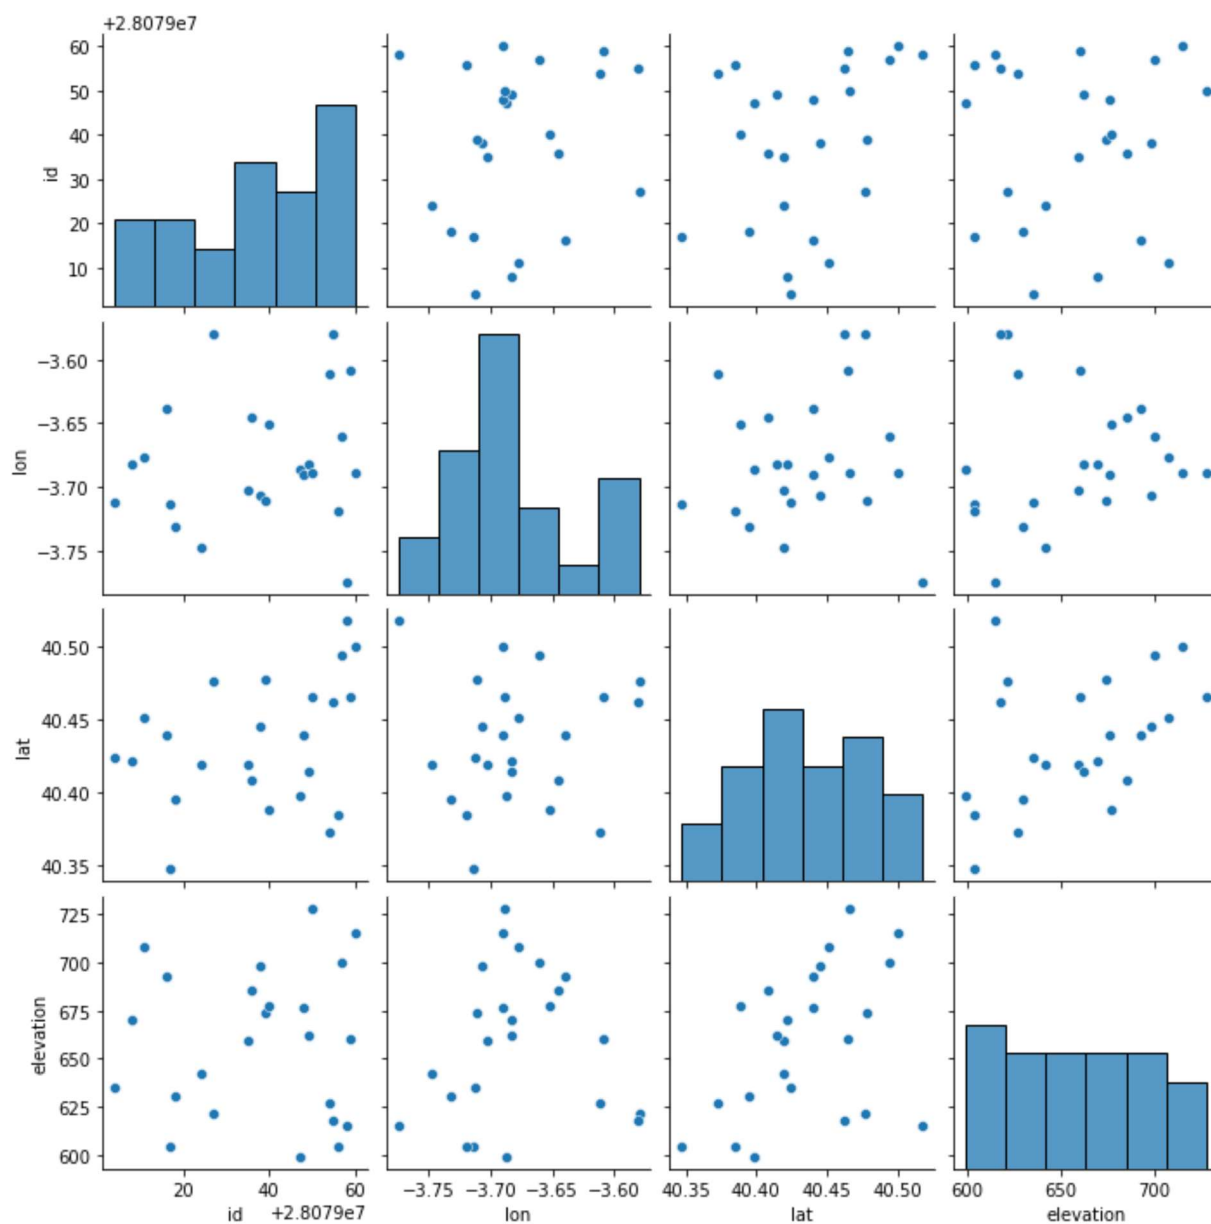
```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 24 entries, 0 to 23
Data columns (total 6 columns):
#   Column      Non-Null Count  Dtype
---  ---
0   id           24 non-null    int64
1   name         24 non-null    object
2   address      24 non-null    object
3   lon          24 non-null    float64
4   lat          24 non-null    float64
5   elevation    24 non-null    int64
dtypes: float64(2), int64(2), object(2)
memory usage: 1.2+ KB
```

In [8]: df.columns

Out[8]: Index(['id', 'name', 'address', 'lon', 'lat', 'elevation'], dtype='object')

```
In [9]: sns.pairplot(df)
```

```
Out[9]: <seaborn.axisgrid.PairGrid at 0x1f561f42af0>
```




```
In [18]: sns.distplot(data['lat'])
```

C:\ProgramData\Anaconda3\lib\site-packages\seaborn\distributions.py:2557: FutureWarning: `distplot` is a deprecated function and will be removed in a future version. Please adapt your code to use either `displot` (a figure-level function with similar flexibility) or `histplot` (an axes-level function for histograms).

```
warnings.warn(msg, FutureWarning)
```

```
Out[18]: <AxesSubplot:xlabel='lat', ylabel='Density'>
```



MODEL BUILDING

1.Linear Regression

```
In [11]: df1=df[['id', 'lon', 'lat', 'elevation']]
```

```
In [26]: x=df1[['id', 'lon', 'lat']]  
y=df1[['elevation']]
```

```
In [27]: #split the dataset into training and test  
from sklearn.model_selection import train_test_split  
x_train,x_test,y_train,y_test=train_test_split(x,y,test_size=0.3)
```


```
In [28]: from sklearn.linear_model import LinearRegression  
lr=LinearRegression()  
lr.fit(x_train,y_train)
```

Out[28]: LinearRegression()

```
In [29]: print(lr.intercept_)  
  
[12778338.31468202]
```

```
In [30]: prediction = lr.predict(x_test)  
plt.scatter(y_test,prediction)
```

Out[30]: <matplotlib.collections.PathCollection at 0x1f564ac38e0>



```
In [31]: print(lr.score(x_test,y_test))  
  
0.033234207960623596
```

2.Ridge Regression

```
In [32]: from sklearn.linear_model import Ridge
```

```
In [33]: rr=Ridge(alpha=10)
rr.fit(x_train,y_train)
```

```
Out[33]: Ridge(alpha=10)
```

```
In [34]: rr.score(x_test,y_test)
```

```
Out[34]: -0.1321146337718111
```

3.Lasso Regression

```
In [35]: from sklearn.linear_model import Lasso
```

```
In [36]: la=Lasso(alpha=10)
la.fit(x_train,y_train)
```

```
Out[36]: Lasso(alpha=10)
```

```
In [37]: la.score(x_test,y_test)
```

```
Out[37]: -0.1277802223297897
```

4.ElasticNet Regression

```
In [38]: from sklearn.linear_model import ElasticNet
en=ElasticNet()
en.fit(x_train,y_train)
```

```
Out[38]: ElasticNet()
```

```
In [39]: print(en.coef_)
```

```
[0.09020638 0.          0.55137415]
```

```
In [40]: print(en.predict(x_test))
```

```
[653.66972253 657.56249315 655.37215378 656.63241645 657.06056021
 652.57831674 655.65729398 656.55624407]
```

```
In [41]: print(en.score(x_test,y_test))
```

```
-0.13260909664647325
```

5.Logistic Regression

```
In [42]: from sklearn.linear_model import LogisticRegression
```

```
In [49]: feature_matrix = df1.iloc[:,0:6]
target_vector = df1.iloc[:, -1]
```

```
In [50]: feature_matrix.shape
```

```
Out[50]: (24, 4)
```

```
In [51]: target_vector.shape
```

```
Out[51]: (24,)
```

```
In [52]: from sklearn.preprocessing import StandardScaler
```

```
In [53]: fs=StandardScaler().fit_transform(feature_matrix)
```

```
In [54]: logr = LogisticRegression()
logr.fit(fs,target_vector)
```

```
Out[54]: LogisticRegression()
```

```
In [57]: observation=[[1,2,3,4]]
```

```
In [58]: prediction=logr.predict(observation)
print(prediction)
```

```
[715]
```

```
In [59]: logr.classes_
```

```
Out[59]: array([599, 604, 615, 618, 621, 627, 630, 635, 642, 659, 660, 662, 670,
        674, 676, 677, 685, 693, 698, 700, 708, 715, 728], dtype=int64)
```

```
In [60]: logr.score(fs,target_vector)
```

```
Out[60]: 0.8333333333333334
```

6.Random Forest

```
In [63]: df1=df[['id', 'lon', 'lat', 'elevation']]
x=df1[['id', 'lon', 'lat']]
y=df1[['elevation']]
```



```
In [64]: from sklearn.model_selection import train_test_split
x_train,x_test,y_train,y_test=train_test_split(x,y,test_size=0.3)
```

```
In [65]: from sklearn.ensemble import RandomForestClassifier
rfc = RandomForestClassifier()
rfc.fit(x_train,y_train)
```

<ipython-input-65-6b9282c2a062>:3: DataConversionWarning: A column-vector y was passed when a 1d array was expected. Please change the shape of y to (n_samples,), for example using ravel().

```
rfc.fit(x_train,y_train)
```

```
Out[65]: RandomForestClassifier()
```

```
In [66]: parameters = {'max_depth':[1,2,3,4,5],
                        'min_samples_leaf':[5,10,15,20,25],
                        'n_estimators':[10,20,30,40,50]}
```

```
In [67]: from sklearn.model_selection import GridSearchCV

grid_search = GridSearchCV(estimator=rfc,param_grid=parameters,cv=2,scoring='acc
grid_search.fit(x_train,y_train)
```

C:\ProgramData\Anaconda3\lib\site-packages\sklearn\model_selection_split.py:666: UserWarning: The least populated class in y has only 1 members, which is less than n_splits=2.

warnings.warn(("The least populated class in y has only %d"

C:\ProgramData\Anaconda3\lib\site-packages\sklearn\model_selection_validation.py:593: DataConversionWarning: A column-vector y was passed when a 1d array was expected. Please change the shape of y to (n_samples,), for example using ravel().

estimator.fit(X_train, y_train, **fit_params)

C:\ProgramData\Anaconda3\lib\site-packages\sklearn\model_selection_validation.py:593: DataConversionWarning: A column-vector y was passed when a 1d array was expected. Please change the shape of y to (n_samples,), for example using ravel().

estimator.fit(X_train, y_train, **fit_params)

C:\ProgramData\Anaconda3\lib\site-packages\sklearn\model_selection_validation.py:593: DataConversionWarning: A column-vector y was passed when a 1d array was expected. Please change the shape of y to (n_samples,), for example using ravel().

estimator.fit(X_train, y_train, **fit_params)

C:\ProgramData\Anaconda3\lib\site-packages\sklearn\model_selection_validation.py:593: DataConversionWarning: A column-vector y was passed when a 1d array was expected. Please change the shape of y to (n_samples,), for example using ravel().

```
In [68]: grid_search.best_score_
```

```
Out[68]: 0.125
```

```
In [69]: rfc_best = grid_search.best_estimator_
```

```
In [70]: from sklearn.tree import plot_tree

plt.figure(figsize=(80,40))
plot_tree(rfc_best.estimators_[5],feature_names=x.columns,filled=True)
```

```
Out[70]: [Text(2232.0, 1087.2, 'gini = 0.836\nsamples = 9\nvalue = [3, 4, 1, 1, 0, 0, 0, 2, 1, 3, 0, 1, 0, 0, 0, 0]')]
```

```
gini = 0.836
samples = 9
value = [3, 4, 1, 1, 0, 0, 0, 2, 1, 3, 0, 1, 0, 0, 0, 0]
```

Results

```
1.Linear regression : 0.033234207960623596
2.Ridge regression : -0.1321146337718111
3.Lasso regression : -0.1321146337718111
4.Elasticnet regression : -0.13260909664647325
5.Logistic regresssion : 0.8333333333333334
6.Random forest regression : 0.125

Hence Logistic regression gives high accuracy for this model.
```